

Systems biology

Fundamental generative structure is invariant across genome-scale reconstructions of *Escherichia coli*

Tomas Veloz^{1,*}

¹Departamento de Matemáticas, Universidad Tecnológica Metropolitana, Santiago 833038, Chile

*To whom correspondence should be addressed. tomas.velozg@utem.cl

Draft prepared for internal review prior to submission.

Abstract

Motivation: Chemical Organization Theory (COT) identifies the persistent (closed, self-maintaining) subnetworks of a reaction network, but exhaustive computation of these organizations is combinatorially intractable beyond a few hundred reactions. A companion theoretical result shows that *elementary semi-organizations* (ESOs) – the productively novel, irreducible building blocks of every organization – can be enumerated directly from a reaction network’s elementary reaction closures (ERCs) via fundamental synergy and complementarity relations, without exhaustive search. Whether this generative structure is stable under genome-scale reconstruction curation, or reshaped by environment the way it is in small hand-built models, has not previously been tested.

Results: We compute ESOs for four successive genome-scale reconstructions of *Escherichia coli* K-12 MG1655 (E_COLI_CORE, iAF1260, iJO1366, iML1515; 2000–2017) under matched environmental (inflow) regimes. ESO size is essentially invariant across the three genome-scale reconstructions (23–31 species) despite 15 years of curation and an ~11% growth in elementary reaction closures per generation; ESO count grows instead, consistent with a conserved metabolic core accumulating parallel, redundant routes as annotation improves. Supplying each reconstruction’s full native medium (trace metals and cofactor precursors absent from the core model) roughly triples mean ESO size, consistent with cofactor-dependent respiratory enzymes bridging otherwise separate modules. Linear-programming verification of self-maintenance – confirmed by code inspection to operate purely on already-enumerated ERC-combination closures rather than by species-level search, and cheap even at genome scale (<0.1 s/ESO) – shows only a minority of ESOs are genuine *elementary organizations* (EOs, the elementary case of Centler *et al.*’s 2008 organization concept): 29–40% for the hand-built core model but a stable 11–15% across all three genome-scale reconstructions and every environment tested, a threefold drop invisible to any core-model-only study. Extending the comparison to three further organisms (a cyanobacterium, a Gram-positive bacterium and a eukaryote) under the same matched environments shows this reconstruction-curation invariance is organism-specific – ESO count and size both vary substantially by organism – while two directional effects, carbon-starvation shrinkage and (in three of four organisms tested, most strongly in an anaerobic metal-reducer) the trace-metal fusion effect, replicate across every organism tested. Testing a specific dual-ancestry hypothesis for the eukaryote in that comparison (*S. cerevisiae*: cytoplasmic/nuclear machinery of archaeal origin operating alongside mitochondrial machinery of bacterial, endosymbiotic origin), using its reconstruction’s own compartment tags, we find no compartment-pure ESO of either kind under any medium tested – 100% of ESOs already draw jointly on both compartments – consistent with complete, long-since-fused integration of the two lineages rather than two separable networks operating in parallel. Reproducing the organization tables of Centler *et al.* (2006) exactly, we further show that only 1–3 of their 4 published organizations per scenario are *fundamental* in this sense; the remainder are combinatorial artefacts of unconsumed environmental species.

Availability and implementation: All code, converted networks and result tables are available in the pyCOT repository (module `pyCOT.analysis.organizations`); this manuscript’s source, figures and data are in `projects/COT_Fundamental_Generators_Exploration/paper_ecoli_comparative`.

Contact: tomas.velozg@utem.cl

Supplementary information: Full per-scenario tables for the reconstruction and organism comparisons are available as `results.csv` in `outputs/ecoli_epm_comparison` and `outputs/cross_organism_epm_comparison` respectively, alongside this manuscript.

1 Introduction

Chemical Organization Theory (COT) explains the persistence of a reaction network through its *organizations*: species sets that are closed (no novel species is produced) and self-maintaining (a non-negative flux exists that keeps every triggered reaction active while never depleting a species in the set) (Dittrich and Speroni di Fenizio, 2007). Every fixed point of a dynamical system built on a reaction network corresponds to an organization, making the organization lattice a structural, simulation-free map of where a network’s stable behaviours can live. The theory’s practical reach, however, has been limited from the outset: enumerating organizations by combining closed sets is combinatorially explosive, and the theory’s own founders flagged efficient computation as an open problem (Dittrich and Speroni di Fenizio, 2007; Centler *et al.*, 2008). Centler *et al.* (2008) answered this with three exhaustive/heuristic algorithms – constructive, flux-based, and a heuristic subset-finder for large models – that remain, to our knowledge, the reference computational approach to COT; they handled *E. coli*-scale central metabolism (~600 reactions) only with parallel computation. Modern genome-scale reconstructions exceed 2 400 reactions (Orth *et al.*, 2011), and human metabolism reconstructions exceed 8 000, placing exhaustive organization search of this kind well out of reach; the generative approach we use here (Section 2) sidesteps that search rather than parallelising it.

COT sits within a broader family of formalisms for extracting persistent or self-sustaining structure from a reaction network without simulating its dynamics. Reflexively autocatalytic food-generated (RAF) sets (Steel *et al.*, 2013) ask a structurally related question – which subnetworks can be built up from food alone with every reaction catalysed by something already present – an origin-of-life-motivated research programme in its own right (Hordijk and Steel, 2018) – and Steel *et al.* (2013) explicitly study the relationship between RAF theory’s *irreducible* RAFs (irrRAFs, the smallest building blocks a RAF can be decomposed into) and COT; irrRAFs and ESOs are, informally, analogous minimal-building-block objects defined from different starting axioms (catalysis vs. closure/self-maintenance), and we return to this connection in Section 4. Integer-hyperflow pathway models (Andersen *et al.*, 2019) instead formalise individual *pathways* as mechanistic, integer-valued flows through a network – a finer-grained, single-route object than an organization or ESO, useful for questions COT does not address (is a specific route chemically realisable) but not for the network-wide persistence question this paper asks; evolutionary-simulation frameworks for metabolic network innovation (Flamm *et al.*, 2010) sit further still from this paper’s scope, asking how such structure *arises* over evolutionary time rather than characterising it directly, as we do here. Tenazinha and Vinga (2011) survey COT alongside these and other integrative modelling frameworks (flux-balance analysis, Petri nets, hybrid and stochastic models), situating it as one of comparatively few approaches that characterise a network’s structural, simulation-free possibility space rather than a single optimal or sampled trajectory through it.

A companion theoretical contribution (Veloz and Bassi, 2026) resolves this by identifying the *generative* structure underlying organizations directly: every semi-organization can be built by combining *elementary reaction closures* (ERCs) through *fundamental synergy* (a combination that activates a reaction neither part could trigger alone) and *fundamental complementarity* (one part supplies a species the other requires), without ever searching the much larger space of generatively irrelevant combinations – adding an inert species, or joining two unrelated closed sets, both produce a “new” semi-organization whose status was already determined by its parts. Restricting to fundamental combinations yields the *elementary semi-organizations* (ESOs): the productively novel, irreducible cores

from which the full organization lattice is built. That paper validated the approach statistically across 438 networks, showing ERC count and fundamental-relation density scale favourably with network size, but it did not ask a longitudinal or comparative biological question: does this generative structure change as a single organism’s metabolic reconstruction is refined over time, or as its environment changes, at genome scale?

The second question already has a celebrated small-scale answer. Centler *et al.* (2006) computed the full organization lattice of a 92-species, 168–171-reaction hand-built model of *E. coli* sugar metabolism under five inflow regimes (environments), and showed each environment selects a different set of organizations from the same underlying network – inflow, in COT, stands in for environment. That result used exhaustive, pre-generative-theory methods and a model an order of magnitude smaller than even the smallest genome-scale BiGG reconstruction of the same organism. Whether the same qualitative environment-sensitivity holds at genome scale, and whether reconstruction curation itself reshapes generative structure, are open, to our knowledge previously unaddressed questions – a literature search for chemical- organization or generative analyses applied comparatively across the iAF1260/iJO1366/iML1515 reconstruction lineage returned no prior work.

This paper reports ESO-level results addressing both questions for four successive genome-scale reconstructions of *E. coli* K-12 MG1655, validates the pipeline by exactly reproducing Centler *et al.*’s own published organizations (decomposed for the first time into their fundamental and spurious, non-fundamental, parts), and then asks a third question the reconstruction-lineage design cannot: is the resulting picture specific to *E. coli*, or does it hold across organisms? We take a first step on that question with four further genome-scale reconstructions spanning a cyanobacterium, a Gram-positive bacterium, a eukaryote, and an anaerobic metal-reducer.

2 Approach

2.1 A worked example: ERCs, fundamental synergy and complementarity

Before the genome-scale results, Fig. 1 works through the vocabulary above on a small, real fragment of *E. coli* central metabolism – six literal BiGG reactions from E_COLI_CORE (PGI, the reverse direction of ADK1, PFK, FBA, TPI, GAPD), with nad_c and pi_c as the only declared food. Two ERCs are not reachable from food alone: closing g6p_c alone through PGI gives $E_1 = \{g6p_c, f6p_c\}$, and closing adp_c alone through ADK1 gives $E_2 = \{adp_c, amp_c, atp_c\}$; neither, alone, satisfies PFK’s requirement for both atp_c and f6p_c together. Their union does: E_1 and E_2 enter a *fundamental synergy*, jointly triggering PFK (and the FBA/TPI cascade it unlocks) to produce a new, larger closure E_6 that neither parent’s closure contains. Downstream, GAPD’s requirement for g3p_c – absent from E_1 , E_2 , or the food closure – is met once E_6 ’s cascade supplies it: a chain of *fundamental complementarities*, each labelled by the single species one ERC supplies to extend another. This six-reaction fragment already exhibits every relation type the genome-scale pipeline searches for; the difference at scale is combinatorial breadth, not a different kind of object.

2.2 Reconstructions and data

We use four reconstructions of *Escherichia coli* str. K-12 substr. MG1655, retrieved from the BiGG Models database (King *et al.*, 2016): the “core” educational model E_COLI_CORE (2000; 95 reactions) (Orth *et al.*, 2010), iAF1260 (2007; 2 382 reactions) (Feist *et al.*, 2007), iJO1366 (2011; 2 583 reactions) (Orth *et al.*, 2011), and iML1515 (2017; 2 712 reactions) (Monk *et al.*, 2017) – successive refinements of the same organism




Fig. 1. A worked example on six real *E. coli* CORE reactions (Section 2.1). (a) The bipartite reaction network: two food species (*nad_c*, *pi_c*) in blue, all other species in grey, reactions as boxes. (b) The resulting ERC hierarchy: grey arrows are containment, the orange diamond is the network’s one fundamental synergy ($E_1 + E_2 \rightarrow E_6$), and blue dashed arrows are the fundamental complementarity chain, each labelled by the species supplied. *E0*, the food closure itself, is a separate persistent ERC not otherwise related to E_1 – E_6 in this small example.

separated by 17 years of curation. Models were converted to pyCOT reaction-network format from the BiGG JSON API. Two data-fidelity issues were identified and corrected during conversion, independently of this study’s results: silently dropped uptake reactions for one unrelated archaeal model (not used here), and scientific-notation stoichiometric coefficients (e.g. biomass-reaction terms such as $5e-05$) being mistokenised by the reaction-string parser into spurious pseudo-species; both are corrected upstream of all results reported here.

For the cross-organism comparison (Section 3.6) we add four further BiGG genome-scale reconstructions, chosen for network size comparable to iAF1260/iJO1366/iML1515 and for organism diversity: *Synechocystis* sp. PCC 6803 (iJN678; 1 071 reactions), *Bacillus subtilis* (iYO844; 1 583 reactions), *Saccharomyces cerevisiae* (iND750; 1 703 reactions), and *Geobacter metallireducens* (iAF987; 1 646 reactions).

2.3 Environment (inflow regime) construction

Following Centler *et al.* (2006), an environment is specified by which species enter the network unconditionally (an inflow reaction with empty support). We define three *core* regimes using only the seven species common to all four reconstructions (*co2_e*, *glc_D_e*, *h2o_e*, *h_e*, *nh4_e*, *o2_e*, *pi_e*): *aerobic* (all seven), *anaerobic* (aerobic minus *o2_e*), and *carbon-starved* (aerobic minus *glc_D_e*). Because these regimes use only species present in every reconstruction, results are directly comparable across reconstruction generations. For the three genome-scale models we additionally run each model’s own *full native medium* – its complete default BiGG exchange set, 22–25 species including trace metals and cofactor precursors (cobalamin, molybdate, tungstate, several transition metals) that *E. coli* CORE does not represent at all – to isolate what the added reconstruction detail itself contributes, within one model. For the four additional organisms (Section 3.6), whose native exchange sets were not hand-curated in advance, we auto-extract

each model’s default medium as every extracellular species with a pure, unconditional “=> species_e” exchange reaction in its raw BiGG file – a convention we first verified reproduces the four *E. coli* models’ hand-curated food sets exactly. Three of the four organisms (*Synechocystis*, *B. subtilis*, *S. cerevisiae*) support the shared 7-species core regimes; *Geobacter metallireducens* has no *glc_D_e* or *o2_e* exchange at all (an obligate anaerobe that does not use glucose) and is run under its own full native medium only.

2.4 Computational pipeline and scope

For each (reconstruction, regime) pair we compute elementary reaction closures (ERCs; via the species-level closure operator, applied exactly once per reaction’s support), the ERC containment hierarchy, fundamental synergies and complementarities, and elementary semi-organizations (ESOs), using the fundamental-generator algorithm of Veloz and Bassi (2026) (module `pyCOT.analysis.organizations`). We confirmed by direct inspection of this implementation that ESO search itself never recomputes a species-level closure: once ERCs are built, new closures are assembled purely by combining ERCs’ precomputed species sets along fundamental synergy/complementarity edges – a bitmask/Horn-propagation step over the ERC hierarchy (module `fundamental_graph.py`, function `erc_syn_close`) – which is what makes genome-scale search tractable at all; species-level closure computation is confined to the one-time, per-reaction ERC-construction step (module `erc.py`). We additionally verify, via linear programming, which ESOs are self-maintaining, i.e. genuine *elementary organizations* (EOs) in the sense of Centler *et al.* (2008) (Section 3.5); this LP pass is cheap (well under 0.1 s/ESO even at genome scale) because it is applied to an already-enumerated, elementary/order-0 ESO list rather than to the much larger search space full organization enumeration would explore. A per-stage timing characterisation on iAF1260 (1 471–1 690

ERCs across regimes) showed the full ERC-to-ESO-to-EO pipeline completes in 1–2 minutes. We deliberately stop at this elementary/order-0 level: higher-order elementary semi-persistent modules (ESPMs) and the free-species/latent-join completion needed to enumerate non-elementary organizations remain future work (Section 4) – a distinct and computationally much larger problem than the per-ESO LP self-maintenance check reported here. One regime – the fully isolated network (no inflow at all) – was excluded for the three genome-scale reconstructions: its fundamental-relations graph is $\sim 20\times$ denser than any fed regime’s (91 322 vs. 4 387 fundamental synergies on iAF1260), and ESO search there did not complete in practical time; it remains included for *E_COLI_CORE*, where it is fast.

As a validation step independent of the *E. coli* reconstruction comparison, we also ran the full pipeline, including free-species and latent-join completion to a naive full enumeration (Veloz and Bassi, 2026, §3.2) and linear-programming verification, on Centler *et al.*’s own five scenarios, to confirm exact agreement with their published organization tables and to decompose those published results into fundamental and spurious parts.

3 Results

3.1 Validating the pipeline against a published ground truth

Figure 2 reports, for each of Centler *et al.*’s five inflow scenarios, how many of the four published organizations are fundamental (reachable by combining ERCs through fundamental synergy/complementarity alone) versus spurious (reachable only by adding a species that triggers no further reaction, or joining two organizations that share no fundamental relation). All 20 published species sets, across all five scenarios, are reproduced exactly. Only 1 of 4 organizations is fundamental under starvation or glucose; 2 of 4 under lactose or glycerol; 3 of 4 under all-sugars. The spurious organizations are, in every case, generated by three specific species – external glucose, lactose and glycerol (Glc_{ex}, Lac_{ex}, Gly_{ex}) – that the original model never actively consumes unless their corresponding uptake machinery is induced. Biologically, this is exactly the distinction one would want a generative theory to draw: the *fundamental* organizations are the genuinely regulated metabolic-genetic cores doing something, while the additional *published* organizations differ from them only by whichever external nutrients happen to be present but unused – bookkeeping for idle resource availability, not additional regulatory structure. That a purely structural, combinatorial criterion recovers this exact distinction, on a network it was never tuned against, is a substantive check on both the pipeline’s correctness and the theory’s descriptive relevance.

3.2 ESO structure is size-invariant across reconstruction generations

Figure 3a–b compares ESO count and size across all four reconstructions under the three core regimes common to all of them. Two patterns are immediately visible. First, ESO *count* grows with reconstruction depth: from 15–16 for *E_COLI_CORE*, to 180–190 for iAF1260, 199–208 for iJO1366, and 194–223 for iML1515, roughly tracking each model’s $\sim 10\%$ growth in ERC count over the previous generation (Fig. 4). Second, and more strikingly, ESO *size* does not grow with it: across the three genome-scale reconstructions, ESO size under any given core regime falls in a narrow, essentially unchanged range (19–31 species; e.g. aerobic: 23–31 in all three), despite 1 518 to 1 690 ERCs and hundreds of added reactions separating iAF1260 from iML1515.

We read this as computational evidence for the conserved *bow-tie* (hourglass) architecture of metabolism (Csete and Doyle, 2004; Zhao *et al.*, 2006): a small, evolutionarily conserved core of central carbon and energy metabolism (glycolysis, TCA cycle, pentose phosphate pathway,




Fig. 2. Of the four organizations Centler *et al.* (2006) publish per inflow scenario, only 1–3 are fundamental (reachable via ERC combination through fundamental synergy or complementarity alone); the rest are generated only by adding an unconsumed external nutrient.




Fig. 3. ESO count (a) and size range (b, min–max with mean) across four *E. coli* reconstructions under matched core environments, plus each genome-scale model’s own full native medium. ESO size is stable across reconstruction generations under any fixed core regime; the full native medium roughly triples mean ESO size.

core amino-acid biosynthesis) through which a much larger and more variable set of peripheral input and output pathways is routed. Successive BiGG curation of *E. coli* predominantly extends the periphery – additional transporters, secondary and alternative-carbon-source pathways, corrected gene–protein–reaction associations – rather than rewriting central metabolism; under a minimal core medium, that peripheral growth manifests as more numerous *parallel* elementary routes through a core of unchanged size (isozyme-driven redundancy) rather than as larger elementary modules. To our knowledge, this is the first demonstration of bow-tie-consistent invariance at the level of a network’s *generative* (rather than purely topological) structure, and the first applied specifically to the reconstruction-history axis rather than a cross-species comparison.

3.3 Trace-metal cofactors fuse, rather than multiply, elementary modules

Supplying each genome-scale model’s full native medium in place of the seven-species core – adding cobalamin (vitamin B₁₂), molybdate, tungstate, nickel, selenium and several further transition-metal ions – produces a qualitatively different effect from adding more of the same core nutrients: mean ESO size roughly triples (e.g. iAF1260: 25.7 to 63.7 species)




Fig. 4. ERC and ESO counts against reaction count (aerobic core regime), log-log. Both grow with reconstruction size; ESO size does not (Fig. 3b).

while ESO *count* falls slightly (182 to 164). We interpret this mechanistically rather than merely descriptively. Molybdenum/tungsten-cofactor and nickel-dependent enzymes in *E. coli* are predominantly anaerobic respiratory reductases – nitrate reductase, DMSO/TMAO reductase, formate dehydrogenase (a selenocysteine-dependent enzyme, requiring the selenium input as well), and [NiFe]-hydrogenases – each of which couples an otherwise peripheral electron-acceptor pathway directly into the shared NADH/quinone redox pool at the heart of central metabolism. Making these cofactors available does not create new, independent self-sustaining pathways; it *connects* pathways that were separately elementary under the core medium into single larger elementary modules, exactly the size-up, count-down signature observed. This gives the trace-metal result a specific, falsifiable mechanistic reading rather than only a numerical one, and suggests ESO fusion under added cofactors as a general, computable proxy for identifying which specific micronutrients are structurally load-bearing for a given metabolic network – a question current constraint-based methods do not answer directly.

3.4 Environment reshapes structure at genome scale, as it does in the small model

Across all three genome-scale reconstructions, the anaerobic core regime yields more ESOs than the aerobic core regime (e.g. iML1515: 223 vs. 216), while carbon starvation yields fewer and smaller ones than either (e.g. iML1515: 214 ESOs, size 19–27, vs. 216 ESOs, size 23–31, aerobic). Both directions match well-established *E. coli* physiology: anaerobic mixed-acid fermentation offers several parallel, non-exclusive branch points (pyruvate formate-lyase vs. the pyruvate dehydrogenase branch; multiple fermentation end products), a known redundancy signature of facultative anaerobes, consistent with more numerous independent elementary modules; aerobic respiration instead channels flux through one dominant oxidative pathway. Carbon starvation's contraction is the direct structural counterpart of losing the network's energy and carbon backbone – we note explicitly that this is a steady-state, structural claim, not a statement about the dynamic stringent response or other regulatory starvation programmes, which this framework does not model. That the qualitative environment-sensitivity Centler *et al.* demonstrated on a 92-species model reproduces at genome scale, across three independent reconstructions and a ~30-fold increase in network size, is itself a result: it was not guaranteed, and to our knowledge had not previously been shown.

3.5 Most ESOs are not organizations, and reconstruction curation widens the gap

For each of the 16 (reconstruction, regime) combinations above we additionally verified, via linear programming, which ESOs are self-maintaining – i.e. which are genuine *elementary organizations* (EOs) in the sense of Centler *et al.* (2008), the elementary case of that paper's general *organization* concept. Centler *et al.*'s own algorithms compute organizations by extending species sets; ours instead verifies already-generated ERC-combination closures directly (Section 2.4), and the LP check itself is cheap even at genome scale: verifying every ESO in a given reconstruction/regime combination takes well under 15 s in total (0.04–0.08 s/ESO for E_COLI_CORE, 0.03–0.07 s/ESO for the three genome-scale reconstructions; Table 1), several orders of magnitude cheaper than the free-species/latent-join completion full organization search would require.

The resulting EO fraction is far from saturating, and drops sharply with reconstruction depth. Under the three shared core regimes, E_COLI_CORE has 29–40% of its ESOs self-maintaining (5–6 of 15–17); every genome-scale reconstruction, under every regime including its own full native medium, falls instead in a narrow 11–15% band (22–27 of 164–223 ESOs; Table 1) – a threefold drop that is stable across three independently curated reconstructions (iAF1260, iJO1366, iML1515) and persists under every environment tested, including the trace-metal-rich full native medium that roughly triples mean ESO size (Section 3.3).

We read this as a second, independent piece of evidence – alongside Section 3.2's size-invariance result – that a hand-built, centrally-curated small model like E_COLI_CORE is not simply a scaled-down genome-scale reconstruction: it over-represents, by a factor of roughly three, how often a productively novel elementary building block is already a self-maintaining structure on its own, rather than one that must combine further – into a still-larger ESO, or into a non-elementary organization – before it can persist. Practically, ESO count/size (computable directly from ERCs) and EO fraction (an added LP pass) carry different, complementary information, and a study reporting only the former at core-model scale risks a materially misleading picture of how much of an organism's generative structure is already self-sustaining.

To our knowledge, this ESO→EO attrition rate has not been quantified before. Centler *et al.*'s own organization-computing algorithms (Centler *et al.*, 2008) build organizations directly by adding candidate species, without separately reporting how many closed sets along the way fail self-maintenance; and Dittrich and Speroni di Fenizio's foundational thesis (Speroni di Fenizio, 2007), while rigorously separating semi-organization from organization in general reaction systems, demonstrates its computations mainly on *catalytic flow systems* (every species purely a catalyst, never stoichiometrically consumed, as in replicator-equation-style artificial chemistries), where a proven lemma makes every semi-organization automatically an organization – a structural feature real, stoichiometric metabolic reconstructions like the four studied here do not share, which is exactly why the gap reported here needed checking rather than being assumed closed.

3.6 Fundamental generative structure across organisms

The size-invariance result above is specific to one organism across one reconstruction lineage; whether ESO structure itself – as opposed to its stability under curation – is similar across organisms is a different question. We computed ESOs, under the same matched core regimes, for three further BiGG genome-scale reconstructions of comparable network size to iAF1260/iJO1366/iML1515 but different organisms: *Synechocystis* sp. PCC 6803 (iJN678, a cyanobacterium), *Bacillus subtilis* (iYO844, Gram-positive), and *Saccharomyces cerevisiae* (iND750, a eukaryote); a fourth, *Geobacter metallireducens* (iAF987, an obligate anaerobic metal-reducer with no glucose or oxygen exchange at all), was run under its

Table 1. ESO→EO verification across all (reconstruction, regime) pairs of Section 3.2. EO = ESO that additionally passes the LP self-maintenance check of Centler *et al.* (2008)'s elementary case.

Reconstruction	Regime	n_{ESO}	n_{EO}	EO%
e_coli_core	isolated	17	5	29%
e_coli_core	aerobic	15	6	40%
e_coli_core	anaerobic	16	6	38%
e_coli_core	C-starved	16	6	38%
iAF1260	aerobic	182	25	14%
iAF1260	anaerobic	189	25	13%
iAF1260	C-starved	180	26	14%
iAF1260	full native	164	25	15%
iJO1366	aerobic	201	22	11%
iJO1366	anaerobic	208	22	11%
iJO1366	C-starved	199	23	12%
iJO1366	full native	179	22	12%
iML1515	aerobic	216	26	12%
iML1515	anaerobic	223	26	12%
iML1515	C-starved	214	27	13%
iML1515	full native	194	26	13%

own full native medium only, since it cannot support the shared aerobic/anaerobic core regimes. Figure 5 reports ESO count and size range for all five organisms side by side (*E. coli*'s iML1515 repeated from Fig. 3 as reference).

Unlike the reconstruction-lineage result, ESO structure is *not* invariant across organisms: under the matched aerobic core, ESO count ranges from 63 (*Synechocystis*) to 216 (*E. coli*) at broadly comparable ERC counts (668–1 690), and mean ESO size is systematically larger for the eukaryote (*S. cerevisiae*, 38.9–42.9 across core regimes) than for any of the three bacteria (15.9–25.6) – a domain-level signal, not a curation artefact, since all four organisms' reconstructions were built and curated independently. One directional pattern does replicate across all four organisms without exception: carbon starvation (removing `glc__D_e` from the core) shrinks mean ESO size relative to the aerobic core in every case (*E. coli* 21.6 vs. 25.6; *Synechocystis* 21.0 vs. 24.0; *B. subtilis* 15.9 vs. 19.0; *S. cerevisiae* 39.9 vs. 42.9), suggesting the structural link between carbon-backbone availability and elementary-module size (Section 3.4) is not *E. coli*-specific.

The trace-metal fusion effect (Section 3.3) partially replicates as well. Supplying *Synechocystis*'s full native medium – which, consistent with real cyanobacterial physiology, uses `no3_e` (nitrate) rather than `nh4_e` as nitrogen source, an organism-specific substitution the auto-extracted food set recovers without being told to – reproduces the *E. coli* pattern exactly: ESO count falls (63→49) while mean size more than doubles (24.0→50.2). *B. subtilis* shows the same direction at smaller magnitude (89→85 ESOs; 19.0→31.0 mean size). *Geobacter*'s full native medium – acetate as carbon source plus a wide panel of alternative terminal electron acceptors (Fe(III), nitrate-like and chromate species, sulfate/sulfite, and dinitrogen), consistent with its ecological role as a versatile anaerobic and extracellular respirer – produces the largest ESOs of any organism or scenario in this study (mean 68.2, range 64–83), the same qualitative signature taken further: an organism whose native lifestyle depends on many alternative respiratory cofactors shows the fusion effect most strongly. *S. cerevisiae* is the exception, but for an identifiable, non-biological reason: its auto-extracted native food set has only 6 species and, unlike the other three organisms' native sets, is a subset rather than a superset of the 7-species core (missing `co2_e` and `h_e`) – this BiGG reconstruction does not declare boundary/exchange reactions for the vitamins and trace elements a real yeast culture requires, so this comparison is uninformative for that organism rather than contradictory; ESO count and size both

fall slightly (152→149; 42.9→24.3 mean), consistent with a *reduced*, not enriched, medium.

3.7 Testing a specific hypothesis: does *S. cerevisiae*'s dual evolutionary origin leave separable elementary traces?

A biochemist reviewer of an earlier draft raised a specific, testable prediction: eukaryotes like *S. cerevisiae* have two distinct evolutionary origins operating side by side – a nuclear/cytoplasmic lineage inherited from the archaeal host, and a mitochondrial lineage inherited from the bacterial endosymbiont that became the mitochondrion – so ESOs might separate into compartment-pure cytoplasmic modules and compartment-pure mitochondrial modules, with the eukaryote's systematically larger ESO size (Section 3.6) reflecting these two networks operating in parallel. iND750 already carries BiGG's own compartment tags on every species (cytoplasmic/nuclear `_c/_n/_r/_g/_v/_x`, mitochondrial `_m`), letting us test this directly with no new modelling assumption.

The prediction does not hold, in a specific and informative way: under every medium tested (the shared aerobic core, iND750's own auto-extracted native medium, and a supplemented medium described below), 100% of ESOs are *compartment-hybrid* – every single elementary module draws on both the cytoplasmic/nuclear and mitochondrial compartments jointly; not one compartment-pure ESO of either kind was found (0/152 under the aerobic core, 0/149 native, 0/146 supplemented). Rather than two separable networks running in parallel, at the elementary/order-0 level *S. cerevisiae*'s reconstruction shows complete, unconditional integration of its two evolutionary origins: no elementary building block remains that could exist using only one compartment's machinery. We read the eukaryote's systematically larger ESO size (Section 3.6) as a consequence of, rather than an alternative to, this complete fusion – consistent with, and a real-network completion of, the same fusion-not-multiplication signature already reported here for trace-metal cofactors within a single compartment (Section 3.3) and for freshly constructed, only partially hybridized endosymbiotic mergers in a companion study (Veloz, 2026): where a newly built merger shows a mix of hybrid and non-hybrid elementary modules, roughly 1.5 billion years of eukaryotic evolution have apparently erased every trace of separability entirely.

The reviewer's related methodological point – that iND750's default medium lacks the vitamins and trace elements mitochondrial respiration depends on – is correct, and only partially testable with this reconstruction: iND750 declares boundary exchange reactions for a handful of relevant cofactors (biotin, riboflavin, thiamin, pantothenate, myo-inositol, sulfate, potassium, sodium) but none at all for iron, zinc, manganese, copper, magnesium, calcium, molybdate, cobalt, folate or niacin – this specific reconstruction simply does not represent most of a real yeast culture's trace-metal requirements as boundary species. Supplying the core medium plus every cofactor iND750 *does* declare (15 species total) leaves the compartment-hybrid result unchanged (146/146, 100%) but reproduces the trace-metal fusion signature's magnitude almost exactly: mean ESO size rises from 24.3 (native, 6-species medium) to 63.9 (supplemented), a 2.6-fold increase comparable in magnitude to *E. coli*'s own trace-metal fusion effect (Section 3.3), while EO fraction falls from 73% (native) to 55% (supplemented, matching the aerobic-core value exactly). A definitive test of whether full trace-metal supplementation could ever produce a compartment-pure ESO would require a reconstruction that declares those exchange reactions, which iND750 does not.




Fig. 5. ESO count (a) and size range (b, min–max with mean) across five organisms under matched core environments (*E. coli*'s iML1515 repeated from Fig. 3 as reference; *G. metallireducens* has no glucose/O₂ exchange and is shown under its own full native medium only). ESO structure varies by organism at matched network scale; carbon starvation shrinks mean ESO size in all four organisms tested; the trace-metal fusion effect of Section 3.3 replicates in *Synechocystis*, *B. subtilis* and (most strongly) *G. metallireducens*, but not in *S. cerevisiae*, whose reconstruction's default exchange set is a reduced, not enriched, medium relative to the core (see text).

4 Discussion and Extensions

4.1 Biological extensions

The bow-tie-consistent size invariance of Section 3.2 is a statement about *reconstruction curation within one organism*; Section 3.6's four-organism test shows it is not simultaneously a statement about organisms in general – ESO count and mean size both vary substantially by organism at matched network scale and matched core environment, with the eukaryote (*S. cerevisiae*) systematically larger than the three bacteria tested. Two patterns did generalise across every organism tested: carbon-starvation shrinkage, and (in three of the four organisms where it could be tested) the trace-metal fusion effect, including its strongest expression yet in *Geobacter metallireducens*, whose ecology – respiration through many alternative extracellular electron acceptors – is exactly what the fusion mechanism predicts should produce the largest effect. The organism-labelled catalogue of 108 BiGG models we have already assembled (spanning bacteria, archaea and eukaryotes) makes the natural next step a systematic version of Section 3.6 at that scale: is the reconstruction-curation size-invariance result itself organism-general (does *E. coli*'s lineage finding replicate in, say, a *B. subtilis* or yeast reconstruction lineage), and does eukaryote/bacterium ESO-size separation hold up across many more species on each side, or was ours a two-eukaryote-vs-many-bacteria artefact of sampling only one representative of the former. The trace-metal fusion result itself suggests a more targeted, single-organism extension: systematically toggling individual cofactor classes (Mo/W, Ni, Se, B₁₂) against enzyme-cofactor annotation to map, mechanistically rather than post hoc, which specific micronutrients are structurally load-bearing in a given network – directly relevant to minimal-medium design and to identifying environmentally contingent essential functions. A further biological extension links ESO and (once available) organization membership to *E. coli*'s well-characterised experimental gene-essentiality and synthetic-lethality data sets, testing whether structurally fundamental modules are enriched for essential genes – a directly testable, falsifiable bridge from this structural theory to phenotype. Finally, the anaerobic redundancy signature reported here invites comparison, using the same catalogue, between obligate and facultative anaerobes, as a computational proxy for metabolic robustness that does not require flux data – *Geobacter* and *E. coli* here already represent opposite poles of that axis.

A separate, more theoretical extension follows from Section 2.1's connection to RAF theory. Steel *et al.* (2013) study irreducible RAFs (irrRAFs) – the smallest autocatalytic building blocks a reflexively autocatalytic food-generated set decomposes into – and their relationship to COT, from a catalysis-first rather than closure-first starting point. Formalising the correspondence between irrRAFs and ESOs (when does an irrRAF-decomposition and an ERC/fundamental-synergy decomposition of the same network agree, and when and why do they diverge) would connect two literatures that ask structurally adjacent minimal-building-block questions from different axiomatic starting points, and could in principle transfer results in either direction – RAF theory's polynomial-time irrRAF bounds (Steel *et al.*, 2013) into COT, or the generative ESO/ESPM construction into RAF enumeration.

4.2 Methodological extensions

Completing the pipeline to full organizations requires two engineering steps identified but not yet implemented: (i) higher-order elementary semi-persistent module (ESPM) search at genome scale, whose tractability boundary (the excluded isolated-network regime) is not yet characterised; and (ii) an efficient linear-programming verification step reusing a single stoichiometric matrix per network (as already done for organization decomposition into catalysts, overproduced species and fragile circuits) rather than rebuilding a sub-network per candidate. With both in place, the same reconstruction and environment comparisons reported here for ESOs extend naturally to full organization lattices, Hasse diagrams, and catalyst/overproduction/fragile-circuit decompositions, which we expect to sharpen rather than overturn the qualitative findings reported here.

Acknowledgements

Text.

Funding

Text.

References

- Andersen, J. L. *et al.* (2019). Chemical transformation motifs – modeling pathways as integer hyperflows. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, **16**(2), 510–523.
- Centler, F. *et al.* (2006). Chemical organizations in the central sugar metabolism of *Escherichia coli*. In C. Priami, E. Merelli, P. Gonzalez, and A. Omicini, editors, *Transactions on Computational Systems Biology V*, volume 3810 of *Lecture Notes in Computer Science*, pages 74–92. Springer.
- Centler, F. *et al.* (2008). Computing chemical organizations in biological networks. *Bioinformatics*, **24**(14), 1611–1618.
- Csete, M. and Doyle, J. (2004). Bow ties, metabolism and disease. *Trends in Biotechnology*, **22**(9), 446–450.
- Dittrich, P. and Speroni di Fenizio, P. (2007). Chemical organisation theory. *Bulletin of Mathematical Biology*, **69**(4), 1199–1231.
- Feist, A. M. *et al.* (2007). A genome-scale metabolic reconstruction for *Escherichia coli* K-12 MG1655 that accounts for 1260 orfs and thermodynamic information. *Molecular Systems Biology*, **3**, 121.
- Flamm, C. *et al.* (2010). Evolution of metabolic networks: a computational framework. *Journal of Systems Chemistry*, **1**, 4.
- Hordijk, W. and Steel, M. (2018). Autocatalytic sets and the origin of life. *Entropy*, **20**(4), 291.
- King, Z. A. *et al.* (2016). Bigg models: A platform for integrating, standardizing and sharing genome-scale models. *Nucleic Acids Research*, **44**(D1), D515–D522.
- Monk, J. M. *et al.* (2017). iML1515, a knowledgebase that computes *Escherichia coli* traits. *Nature Biotechnology*, **35**(10), 904–908.
- Orth, J. D. *et al.* (2010). Reconstruction and use of microbial metabolic networks: the core *Escherichia coli* metabolic model as an educational guide. *EcoSal Plus*, **4**(1).
- Orth, J. D. *et al.* (2011). A comprehensive genome-scale reconstruction of *Escherichia coli* metabolism – 2011. *Molecular Systems Biology*, **7**, 535.
- Speroni di Fenizio, P. (2007). *Chemical Organization Theory*. Ph.D. thesis, University of Jena.
- Steel, M. *et al.* (2013). Minimal autocatalytic networks. *Journal of Theoretical Biology*, **332**, 96–107.
- Tenazinha, N. and Vinga, S. (2011). A survey on methods for modeling and analyzing integrated biological networks. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, **8**(4), 943–958.
- Veloz, T. (2026). Endosymbiotic fusion of elementary semi-organizations: a computable structural model of evolutionary complexification. *in preparation, submitted to PNAS*.
- Veloz, T. and Bassi, A. (2026). Synergy and complementarity: The generative basis of chemical organizations. *in preparation*.
- Zhao, J. *et al.* (2006). Hierarchical modularity of nested bow-ties in metabolic networks. *BMC Bioinformatics*, **7**, 386.